

# What is a column family database?

NOSQL CONCEPTS




**Miriam Antona**  
Software engineer

# Column family databases - overview

- Derive from **Google BigTable**
- Store data in **column families**
  - group related data
  - frequently accessed together
- Also called **wide column** databases
- Great when dealing with large volumes of data

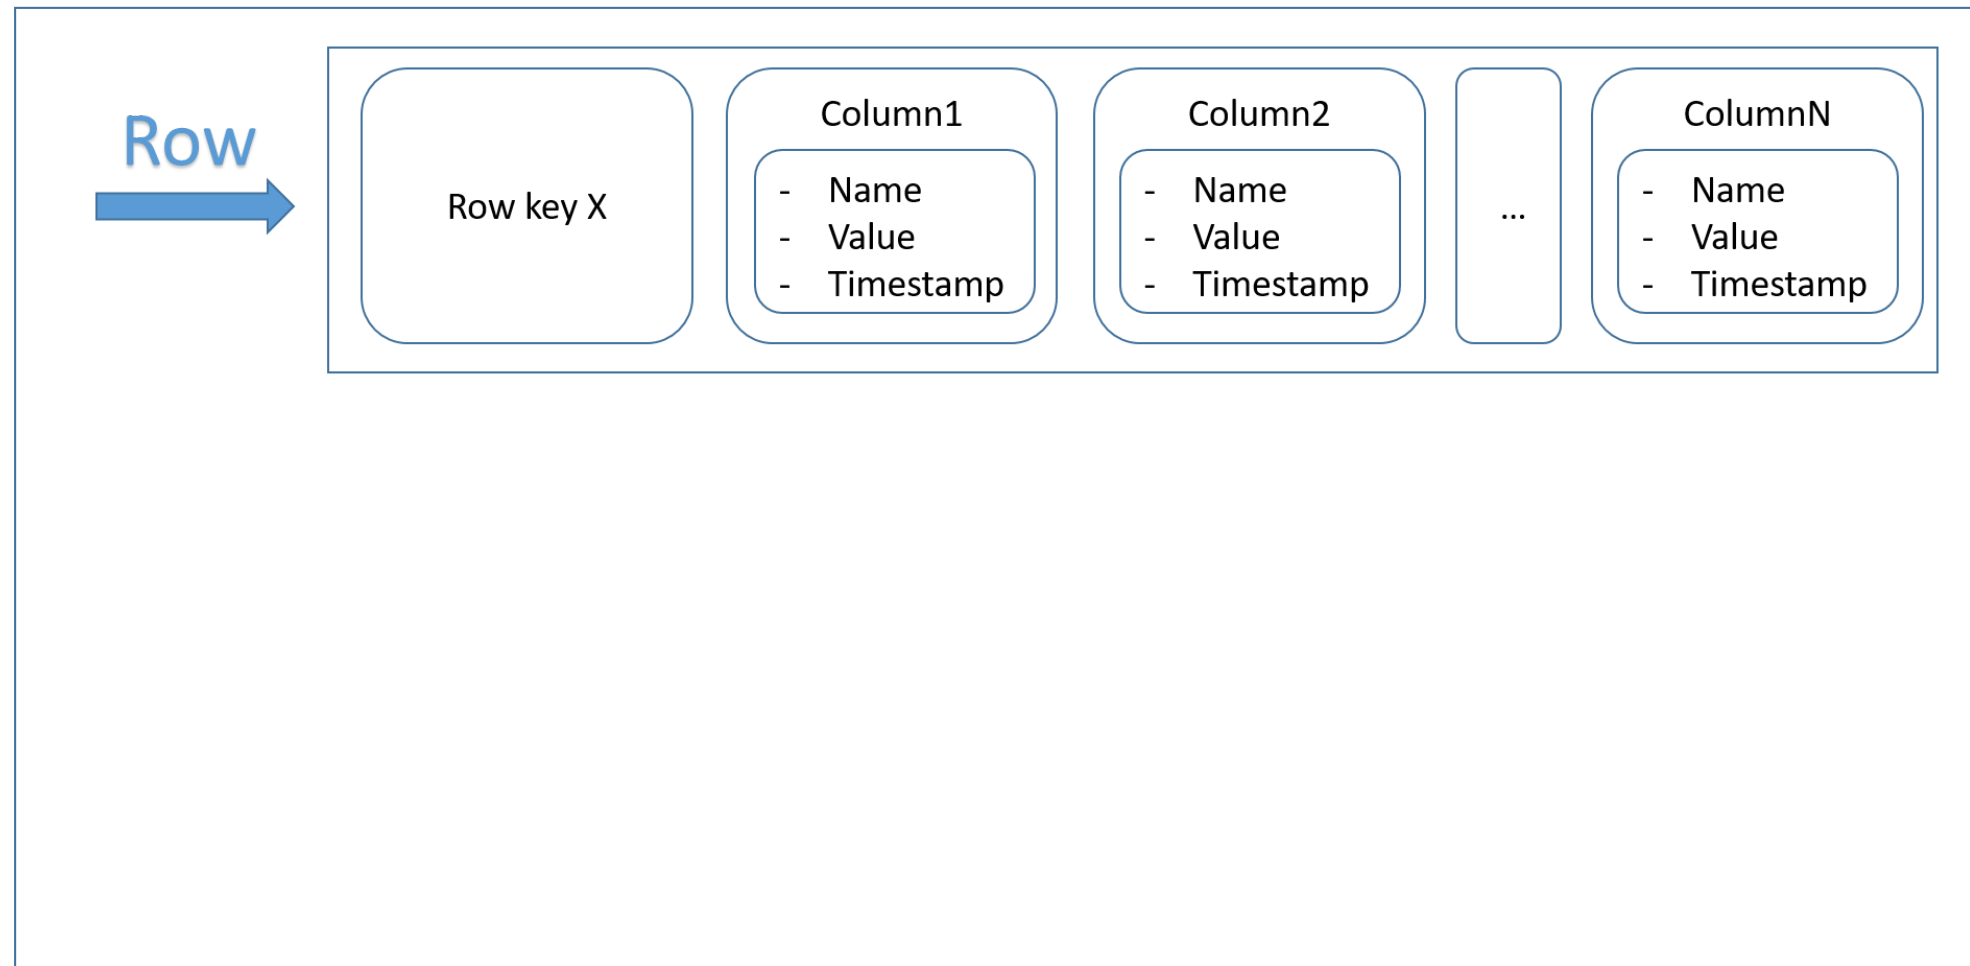
# Column family databases - structure

Column family



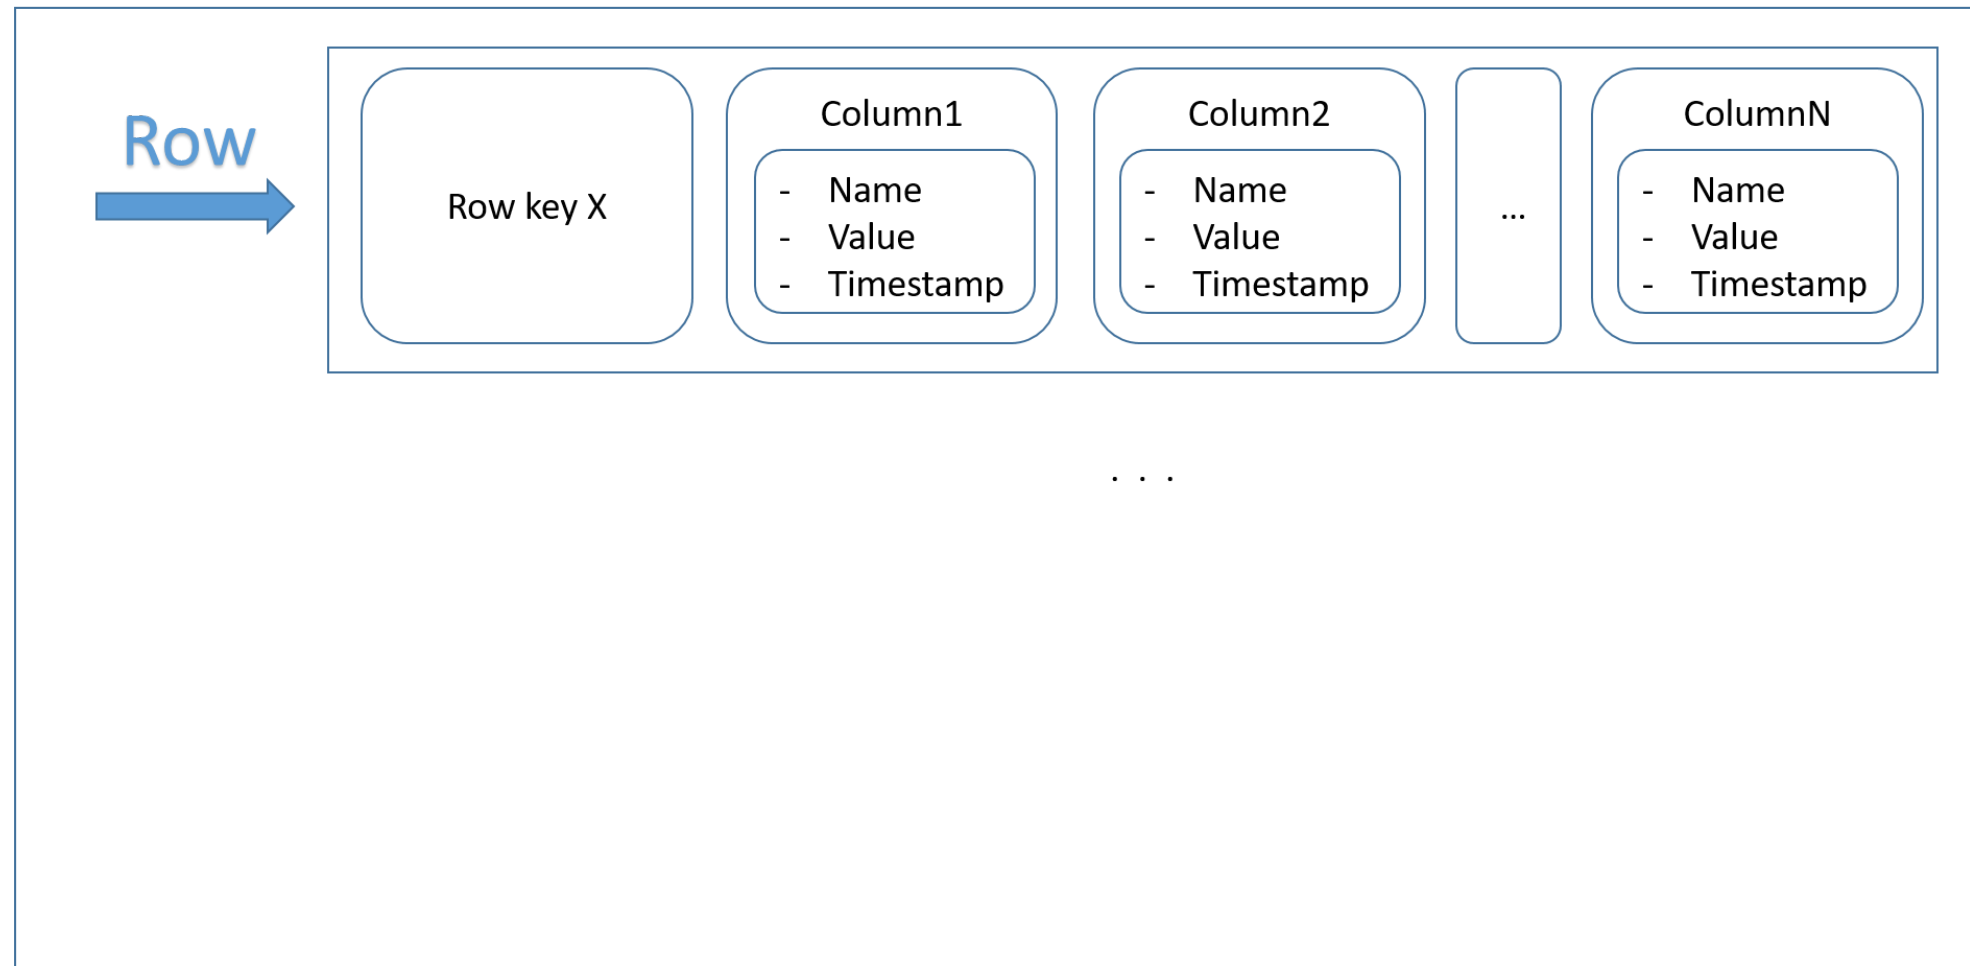
# Column family databases - structure

## Column family



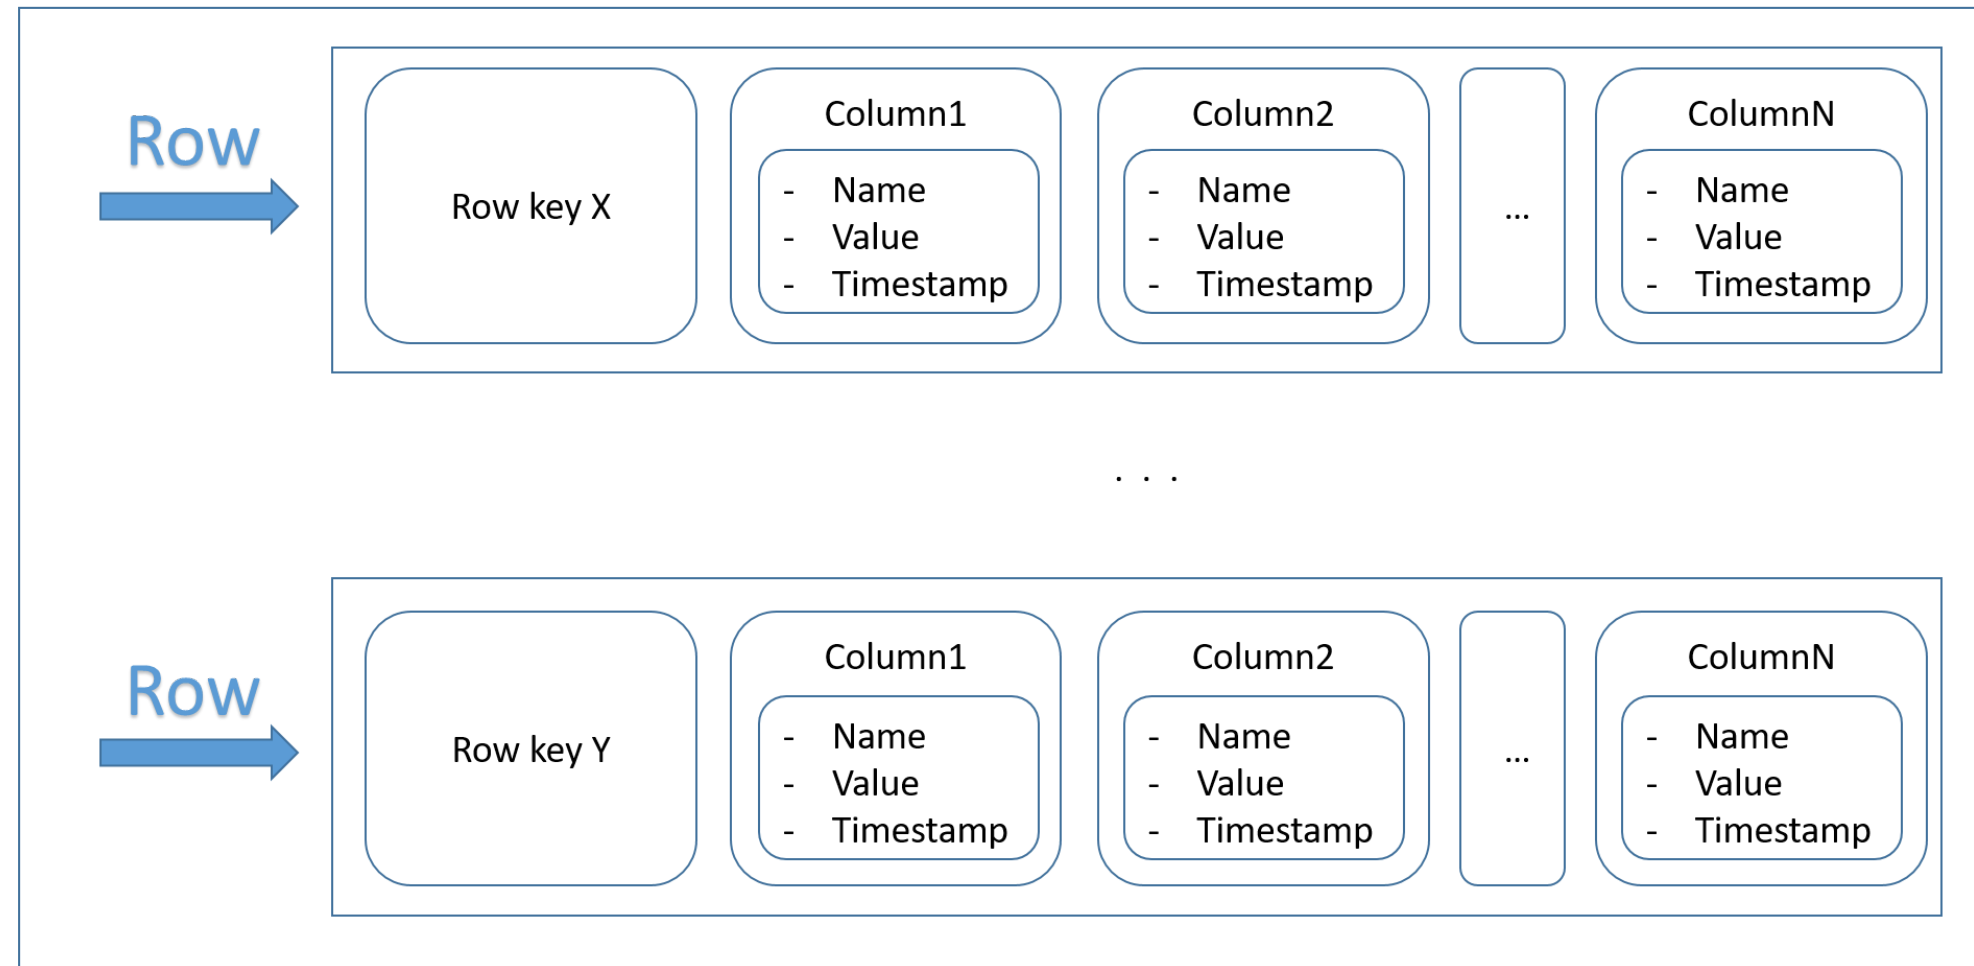
# Column family databases - structure

## Column family



# Column family databases - structure

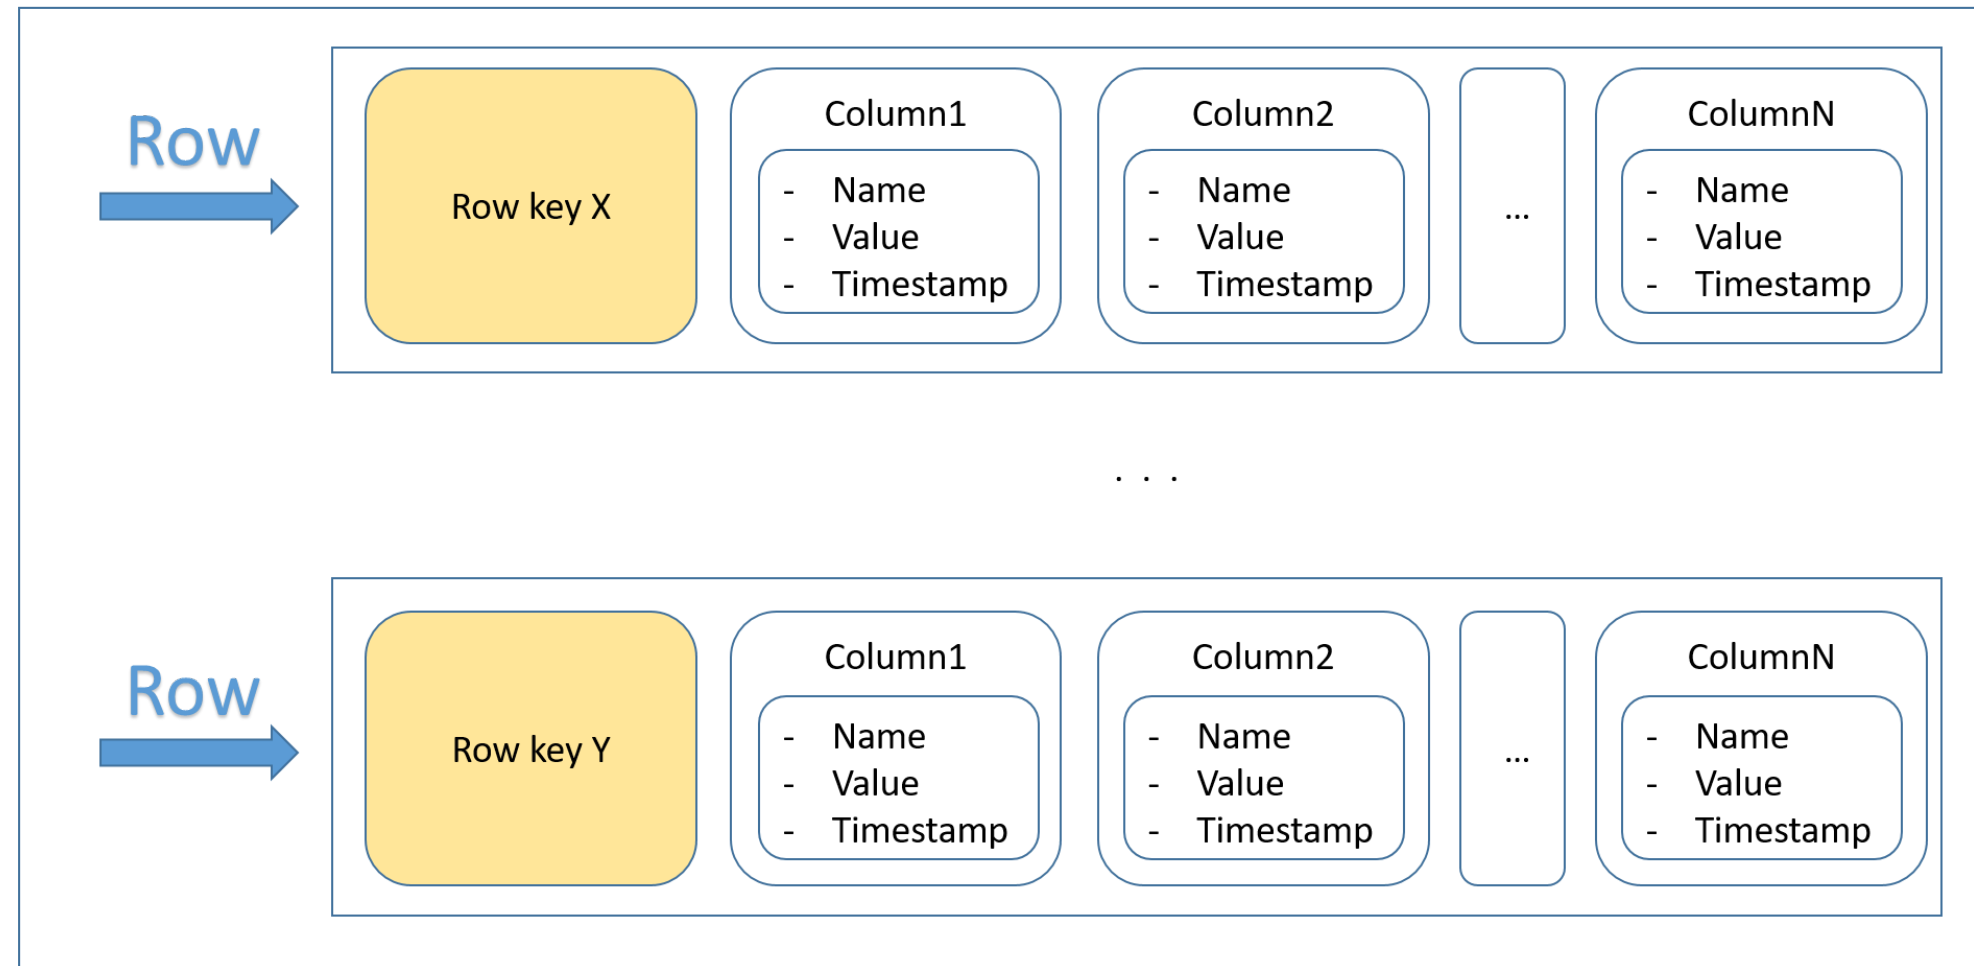
## Column family



- A **column family** is like a **table** in a relational database

# Column family databases - structure

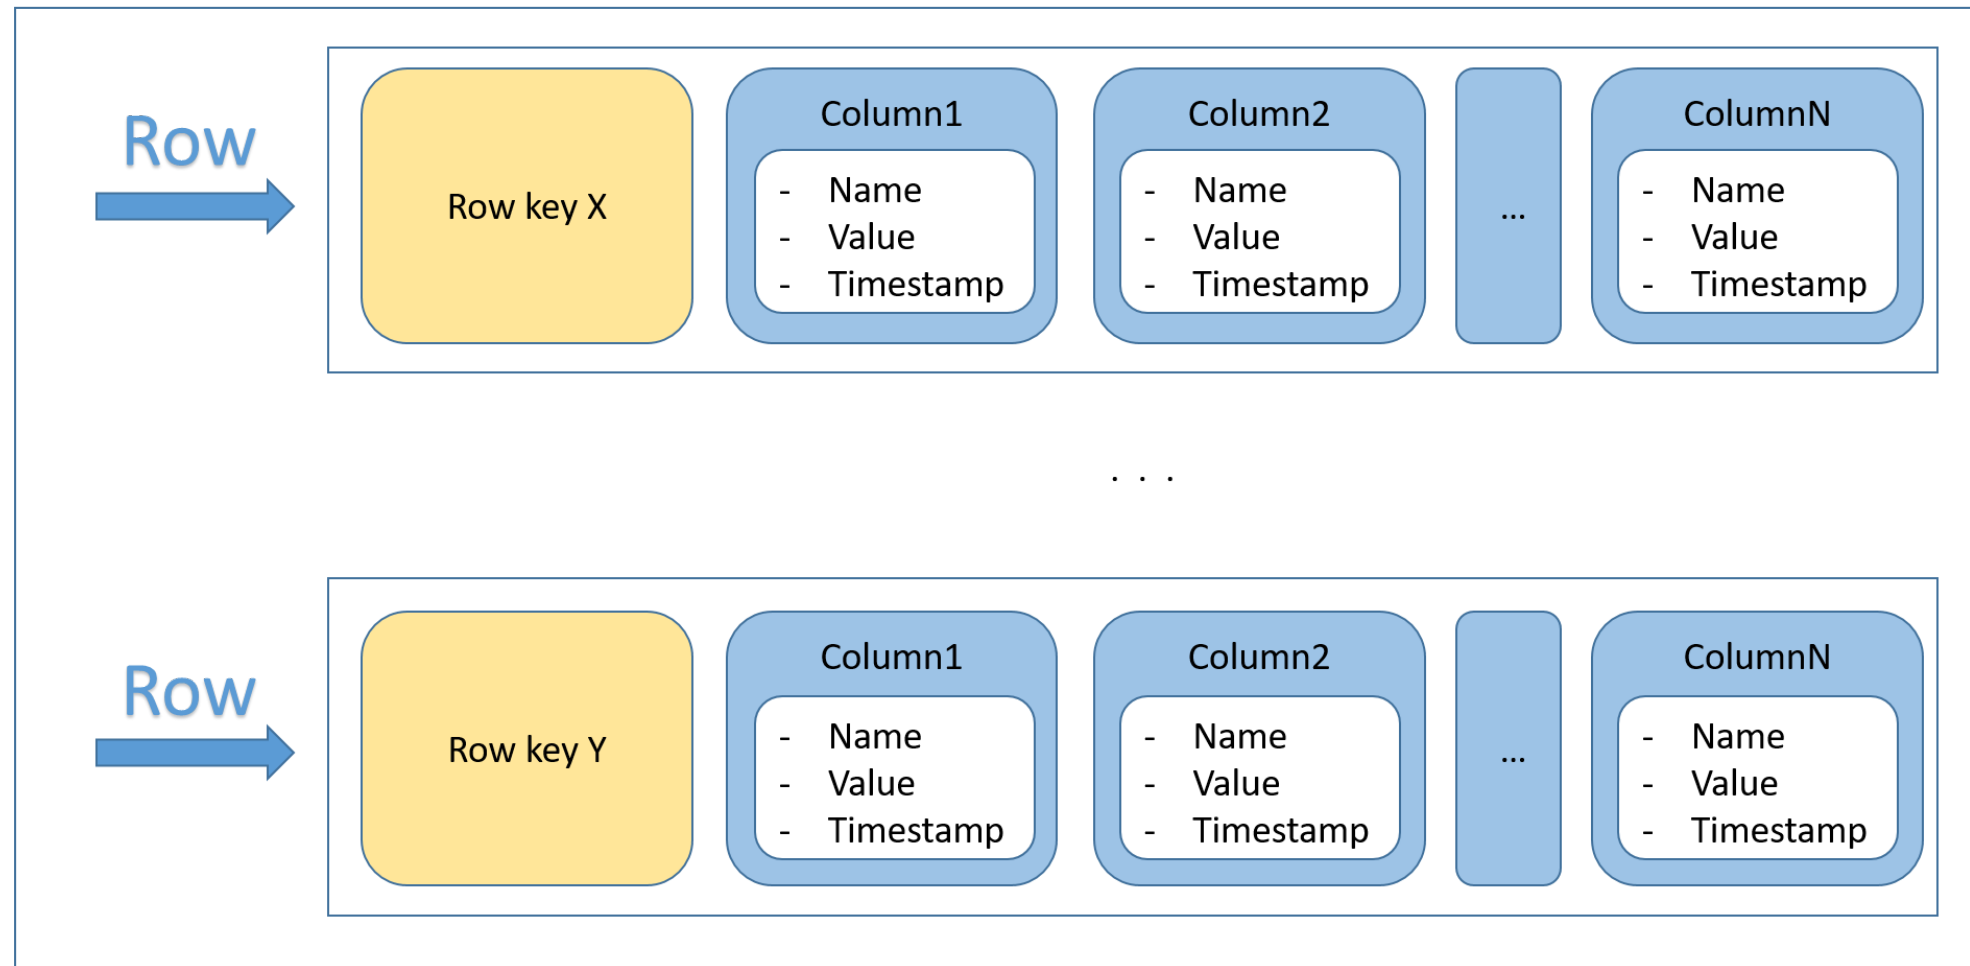
## Column family



- **Row key:** unique identifiers
  - Like primary keys in a relational database

# Column family databases - structure

## Column family




- Each **row** can have different number of columns
  - Columns can be added when needed



# Column family databases - structure


## Column family



- **Parts** of the columns:
  - Name, value, and timestamp

# Column family databases - structure

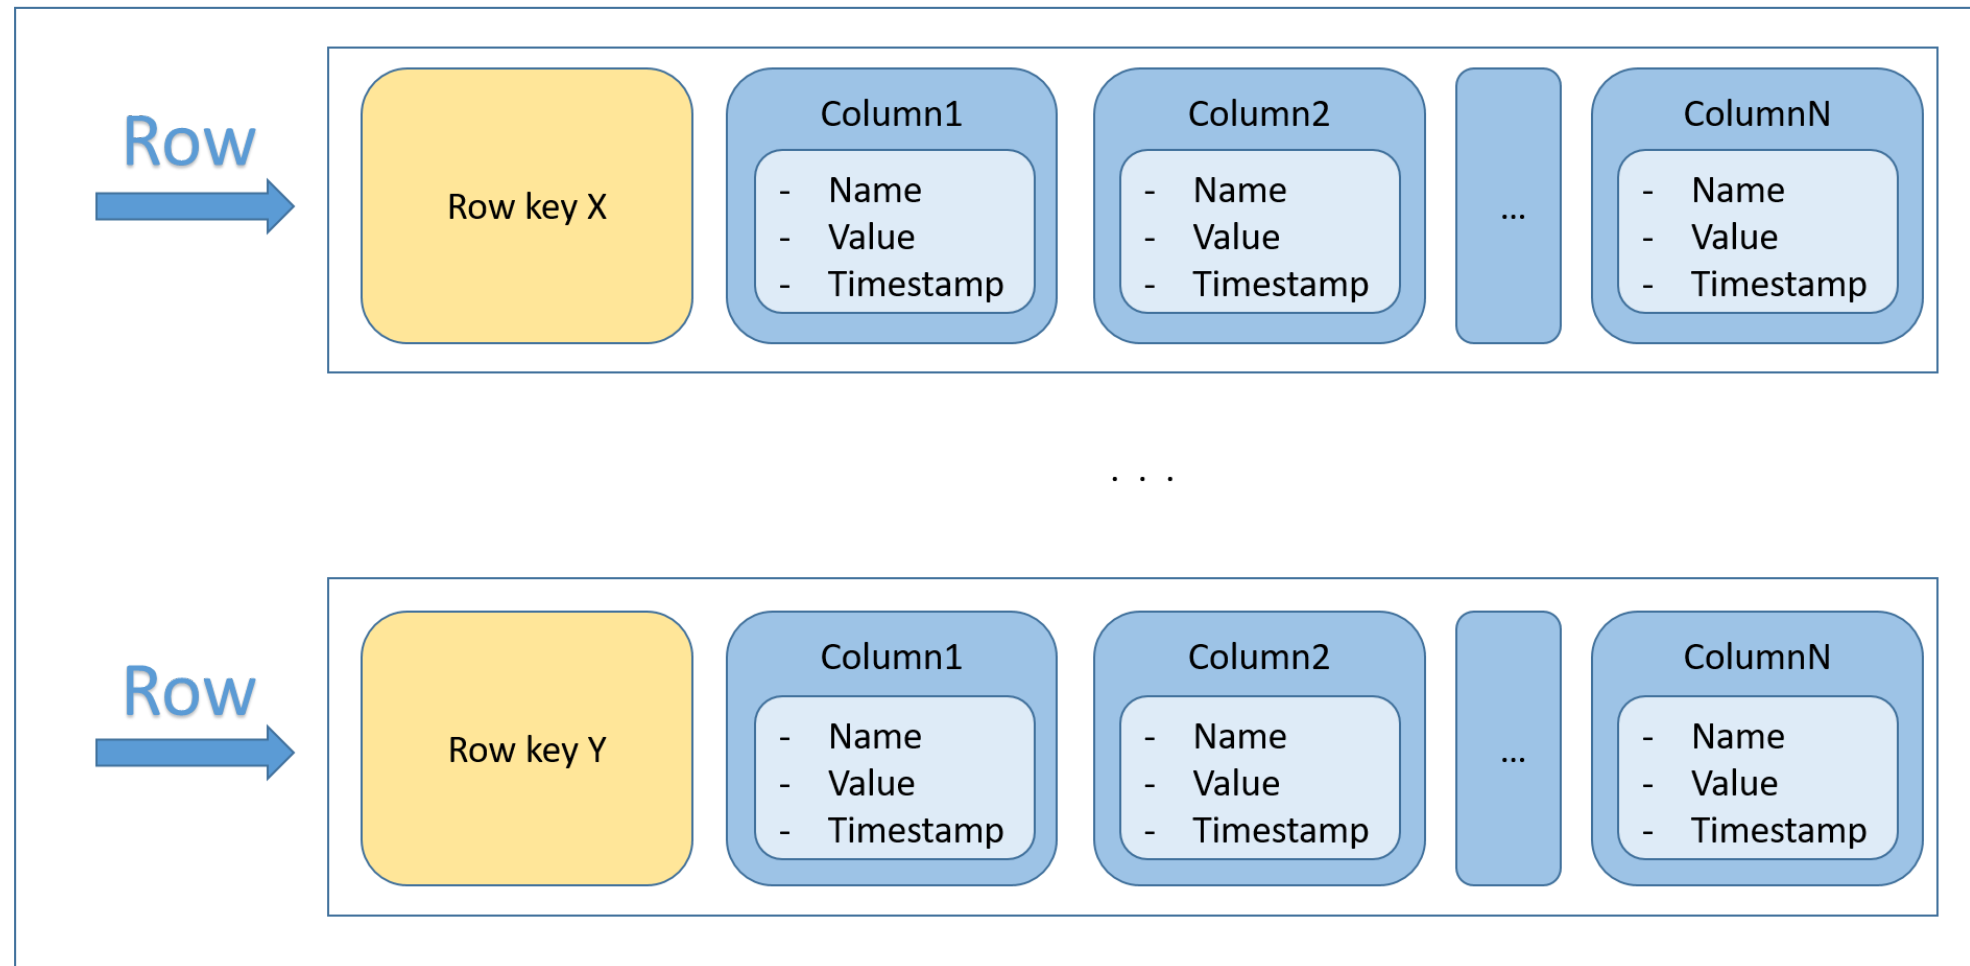
## Column family



- **Value:** specify the type depending on the database

# Column family databases - structure

## Column family



- **Timestamps:** store date and time when the data was inserted.
  - Multiple values of a column

# Column family databases - example

512	first name Carol	last name Harper	email carol@datazy.com	address 123 Sesame St, NY
513	first name Benjamin	last name Lieberman	email ben@datazy.com	
514	first name Peter	last name Stallings	email peter@datazy.com	date of birth 07/04/1984

# Column family databases - designing

- Think about the queries
- No joins
  - Add all the columns we need


# Popular column family databases



**Let's practice!**  
NOSQL CONCEPTS

# Advantages and limitations of column family databases

NOSQL CONCEPTS



**Miriam Antona**  
Software engineer



# Advantages - flexibility

- Rows within a column family can have **different columns**
- **Add new columns** to a row if we need them
- **Avoids** filling with **default values**
- Flexibility mustn't be considered as the only criterion
  - Evaluate key-value and document databases

# Advantages - speed

- Related columns are **stored together on disk**
- **Very fast** writing / retrieving

# Advantages - scalability

- Scale horizontally
  - Sharding across multiple servers

# Advantages - large volumes of data

- Designed to handle **large volumes of data**
  - speed
  - horizontal scalability
  - efficient data compression

# Limitations

- Atomic reads/writes but **no multirow transactions**
- **No joins** support
- **No subqueries** support
- Need to **define the queries** quite well
  - Queries change -> may need to change the column families
  - Can be costly

**Let's practice!**  
NOSQL CONCEPTS

# When to use column family databases

NOSQL CONCEPTS



**Miriam Antona**  
Software engineer


# Suitable cases - general cases

- Large volumes of data
- Extreme write speeds



# Suitable cases - event logging

- Types of events:
  - User logging
  - Errors
  - ...




# Suitable cases - Content Management Systems

- Comments
- Links
- Tags
- ...



# Suitable cases - time-series data

- Weather
- Traffic
- etc.



# Unsuitable cases

- **Prototyping** and at the **beginning** of a project
  - Need to change the queries very frequently
  - Changing the queries -> may imply changing the design of the column families
  - Costly and may slow down the productivity
- Complex queries and joins
- Not dealing with large amounts of data

**Let's practice!**  
NOSQL CONCEPTS


# Apache Cassandra case study

NOSQL CONCEPTS



**Miriam Antona**  
Software engineer

# Apache Cassandra - overview



Apache

# CASSANDRA

- Popular **column family** database
- Originally developed by **Facebook**
- **Open-source**
- Finally became a project of the **Apache Foundation**

# Apache Cassandra - features

- **Distributed**
  - Data is distributed across the nodes of the cluster
  - Every node plays the same role
  - No master node
- **High availability**
- **No single point of failure**
- **Scales horizontally** by adding nodes
- Cassandra client **drivers**: C#, Java, Python, Scala, etc.



# Apache Cassandra - features

- Cassandra Query Language (aka **CQL**)
  - Query data
  - Similar syntax to SQL
  - Tables (for column families), rows, and columns
  - Differences between CQL and SQL:
    - no joins
    - no foreign keys
    - no subqueries, etc.
    - rows can contain a different number of columns

```
SELECT * FROM users WHERE user_id IN (212, 213, 214);
```

# Apache Cassandra - ecosystem

- **Third-party** Cassandra projects, tools, products, and services
  - Cloud offerings
  - Installation tools
  - Developers' frameworks
  - Connectors
  - etc.

# Apache Cassandra - customers

**NETFLIX**



# Bigmate case study - overview

- Location tracking
- Industrial sensor
- Productivity



# Bigmate case study - problem and solution

- **IoT platform:**
  - Ingests and processes **large volumes** of **different data**
  - **Integrate** IoT sensors, devices, and other platforms
  - Process data in **real-time**
  - Scale and deploy across **multiple locations**
  - Application examples:
    - Thermy -> capture the skin temperature of people
    - Warny -> detects possible collisions
- Tested MySQL, MongoDB, Apache Cassandra, etc.
  - Chose **Apache Cassandra**
  - **Scaled better**

# Bigmate case study - results

- **Millions of operations of concurrent** users
- Display 20,000 real-time data points to a single customer
- Fault tolerance (data replication)

# Bigmate case study - results

- **Millions of operations of concurrent** users
- Display 20,000 real-time data points to a single customer
- Fault tolerance (data replication)

<sup>1</sup> <https://cassandra.apache.org/case-studies/>

**Let's practice!**  
NOSQL CONCEPTS